

Introduction to Bayesian statistical modelling

A course with R, Stan, and brms

Ladislav Nalborczyk (UNICOG, NeuroSpin, CEA, Gif/Yvette, France)



Planning

Course n°01: Introduction to Bayesian inference, Beta-Binomial model

Course n°02: Introduction to brms, linear regression

Course n°03: Markov Chain Monte Carlo, generalised linear model

Course n°04: Multilevel models, cognitive models



Setting up a common language

$$y_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta x_i$$

$$\alpha \sim \text{Normal}(60, 10)$$

$$\beta \sim \text{Normal}(0, 10)$$

$$\sigma \sim \text{HalfCauchy}(0, 1)$$

Aim of this session: to understand this type of model.

The components of our models will always be the same and we will always follow the following three steps:

- Build the model (likelihood + priors).
- Update with data to calculate the posterior distribution.
- Interpret the model's estimates and evaluate its predictions. If necessary, modify the model.



A first model



A first model

```
1 library(tidyverse)
2 library(imsb)
3
4 d <- open_data(howell)
5 str(d)
```

```
'data.frame':  544 obs. of  4 variables:
 $ height: num  152 140 137 157 145 ...
 $ weight: num  47.8 36.5 31.9 53 41.3 ...
 $ age   : num  63 63 65 41 51 35 32 27 19 54 ...
 $ male  : int   1 0 0 1 0 1 0 1 0 1 ...
```


```
1 d2 <- d %>% filter(age >= 18)
2 head(d2)
```

	height	weight	age	male
1	151.765	47.82561	63	1
2	139.700	36.48581	63	0
3	136.525	31.86484	65	0
4	156.845	53.04191	41	1
5	145.415	41.27687	51	0
6	163.830	62.99259	35	1

A first model

$$h_i \sim \text{Normal}(\mu, \sigma)$$


```
1 d2 %>%  
2   ggplot(aes(x = height) ) +  
3   geom_histogram(aes(y = ..density..), bins = 20, col = "white") +  
4   stat_function(fun = dnorm, args = list(mean(d2$height), sd(d2$height) ), size = 1)
```



Normal distribution


$$p(x \mid \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[-\frac{1}{2\sigma^2}(\mu - x)^2 \right]$$

```
1 data.frame(value = rnorm(n = 1e4, mean = 10, sd = 1) ) %>% # 10.000 samples from Normal(10, 1)
2   ggplot(aes(x = value) ) +
3   geom_histogram(col = "white")
```



Where does the normal distribution come from?


Some constraints: Some values are highly probable (around the mean μ). The further away they are, the less likely they are (following an exponential decay).



Where does the normal distribution come from?

$$y = \exp \left[-x^2 \right]$$


We extend our function to negative values.



Where does the normal distribution come from?

$$y = \exp \left[-x^2 \right]$$


The inflection points give us a good indication of where most of the values lie (i.e., between the inflection points). The peaks of the derivative show us the inflection points.



Where does the normal distribution come from?

$$y = \exp \left[-\frac{1}{2}x^2 \right]$$


Next, we standardise the distribution so that the two inflection points are located at $x = -1$ and $x = 1$.



Where does the normal distribution come from?

$$y = \exp \left[- \frac{1}{2\sigma^2} x^2 \right]$$


We then add a parameter σ^2 to control the distance between the inflection points.



Where does the normal distribution come from?

$$y = \exp \left[- \frac{1}{2\sigma^2} (\mu - x)^2 \right]$$


We then insert a parameter μ to control the position (central tendency) of the distribution.



Where does the normal distribution come from?

$$y = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[-\frac{1}{2\sigma^2} (\mu - x)^2 \right]$$

But... this distribution does not integrate to 1. We therefore divide by a normalisation constant (the left-hand side), to obtain a proper (valid) probability distribution.



Gaussian model

We are going to build a regression model, but before we add a predictor, let's try to model the distribution of heights.

We want to know which model (distribution) best describes the distribution of heights. We will therefore explore all the possible combinations of μ and σ and rank them by their respective probabilities.


Our aim, once again, is to describe **the posterior distribution**, which will therefore be in some way **a distribution of distributions**.



Gaussian model

We define $p(\mu, \sigma)$, the joint prior distribution of all model parameters. These priors can be specified independently for each parameter, given that $p(\mu, \sigma) = p(\mu)p(\sigma)$.


$$\mu \sim \text{Normal}(174, 20)$$



Gaussian model

We define $p(\mu, \sigma)$, the joint prior distribution of all model parameters. These priors can be specified independently for each parameter, given that $p(\mu, \sigma) = p(\mu)p(\sigma)$.

$$\sigma \sim \text{Uniform}(0, 50)$$




Visualiser le prior

```


1 library(ks)
2 sample_mu <- rnorm(1e4, 174, 20) # prior on mu
3 sample_sigma <- runif(1e4, 0, 50) # prior on sigma
4 prior <- data.frame(cbind(sample_mu, sample_sigma) ) # multivariate prior
5 H.scv <- Hscv(x = prior, verbose = TRUE)
6 fhat_prior <- kde(x = prior, H = H.scv, compute.cont = TRUE)
7 plot(
8   fhat_prior, display = "persp", col = "steelblue", border = NA,
9   xlab = "\nmu", ylab = "\nsigma", zlab = "\n\np(mu, sigma)",
10  shade = 0.8, phi = 30, ticktype = "detailed",
11  cex.lab = 1.2, family = "Helvetica")

```



Prior predictive checking


```
1 sample_mu <- rnorm(1000, 174, 20)
2 sample_sigma <- runif(1000, 0, 50)
3
4 data.frame(x = rnorm(n = 1000, mean = sample_mu, sd = sample_sigma) ) %>%
5   ggplot(aes(x) ) +
6   geom_histogram() +
7   labs(x = "Height (cm)", y = "Number of samples")
```



Likelihood function

```
1 mu_exemple <- 151.23
2 sigma_exemple <- 23.42
3
4 d2$height[34] # exemplary height observation
```

```
[1] 162.8648
```



Likelihood function

We want to compute the probability of observing a certain value of height knowing certain values of μ and σ , that is:

$$p(x \mid \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[-\frac{1}{2\sigma^2}(\mu - x)^2 \right]$$

This **probability density** can be computed using the functions `dnorm`, `dbeta`, `dt`, `dexp`, `dgamma`, etc.

```
1 dnorm(d2$height[34], mu_exemple, sigma_exemple)
[1] 0.01505675
```



Likelihood function

$$p(x | \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[-\frac{1}{2\sigma^2}(\mu - x)^2 \right]$$

Or using a custom function...

```
1 normal_likelihood <- function (x, mu, sigma) {  
2  
3   bell <- exp( (- 1 / (2 * sigma^2) ) * (mu - x)^2 )  
4   norm <- sqrt(2 * pi * sigma^2)  
5  
6   return(bell / norm)  
7  
8 }
```

```
1 normal_likelihood(d2$height[34], mu_exemple, sigma_exemple)
```

```
[1] 0.01505675
```



Posterior distribution

$$p(\mu, \sigma | h) = \frac{\prod_i \text{Normal}(h_i | \mu, \sigma) \text{Normal}(\mu | 174, 20) \text{Uniform}(\sigma | 0, 50)}{\int \int \prod_i \text{Normal}(h_i | \mu, \sigma) \text{Normal}(\mu | 174, 20) \text{Uniform}(\sigma | 0, 50) d\mu d\sigma}$$

$$p(\mu, \sigma | h) \propto \prod_i \text{Normal}(h_i | \mu, \sigma) \text{Normal}(\mu | 174, 20) \text{Uniform}(\sigma | 0, 50)$$

This is the same formula from the previous course, but here considering that there are several observations of height (h_i) and two parameters to be estimated: μ and σ .

To compute the **marginal likelihood** (in green), we therefore need to integrate over two parameters: μ and σ . Here again we realise that the posterior is proportional to the product of the likelihood and the prior.



Posterior distribution - Grid approximation

```
1 # we define a grid of possible values for mu and sigma
2 mu.list <- seq(from = 140, to = 160, length.out = 200)
3 sigma.list <- seq(from = 4, to = 9, length.out = 200)
4
5 # we extend this grid to all possible combinations of mu and sigma
6 post <- expand.grid(mu = mu.list, sigma = sigma.list)
7
8 # we compute the log-likelihood for each observation (under each combination of mu and sigma)
9 post$LL <-
10   sapply(
11     1:nrow(post),
12     function(i) sum(dnorm(
13       d2$height,
14       mean = post$mu[i],
15       sd = post$sigma[i],
16       log = TRUE) )
17   )
18
19 # we compute the (unnormalised) posterior distribution
20 post$prod <-
21   post$LL +
22   dnorm(x = post$mu, mean = 174, sd = 20, log = TRUE) +
23   dunif(x = post$sigma, min = 0, max = 50, log = TRUE)
24
25
```



Posterior distribution - Grid approximation


```
1 # randomly selecting 20 rows from the resulting dataframe
2 post %>% slice_sample(n = 20, replace = FALSE)
```

	mu	sigma	LL	prod	prob
1	155.1759	4.251256	-1418.233	-1426.502	4.612132e-87
2	143.2161	6.211055	-1829.978	-1838.989	3.334659e-266
3	156.6834	6.462312	-1250.548	-1258.749	3.297048e-14
4	155.8794	6.537688	-1237.281	-1245.518	1.837911e-08
5	141.0050	8.221106	-1701.767	-1710.955	1.342455e-210
6	159.0955	8.246231	-1273.181	-1281.286	5.378841e-24
7	142.6131	4.728643	-2471.268	-2480.326	0.000000e+00
8	146.3317	5.507538	-1667.237	-1676.021	1.993230e-195
9	160.0000	6.788945	-1337.375	-1345.446	7.347547e-52
10	143.4171	7.517588	-1608.947	-1617.943	3.329394e-170
11	140.6030	6.437186	-2064.584	-2073.805	0.000000e+00
12	151.6583	5.005025	-1370.982	-1379.432	1.276988e-66
13	156.2814	6.814070	-1236.277	-1244.496	5.108761e-08
14	153.7688	8.849246	-1226.826	-1235.164	5.767419e-04
15	144.1206	7.015075	-1615.495	-1624.437	5.033474e-173
16	143.7186	8.572864	-1506.316	-1515.289	1.271644e-125
17	154.8744	7.517588	-1219.926	-1228.210	6.039964e-01
18	146.3317	8.824121	-1379.469	-1388.252	1.886826e-70
19	140.0000	7.844221	-1828.939	-1838.211	7.264600e-266
20	159.8995	5.180905	-1478.782	-1486.857	2.832570e-113



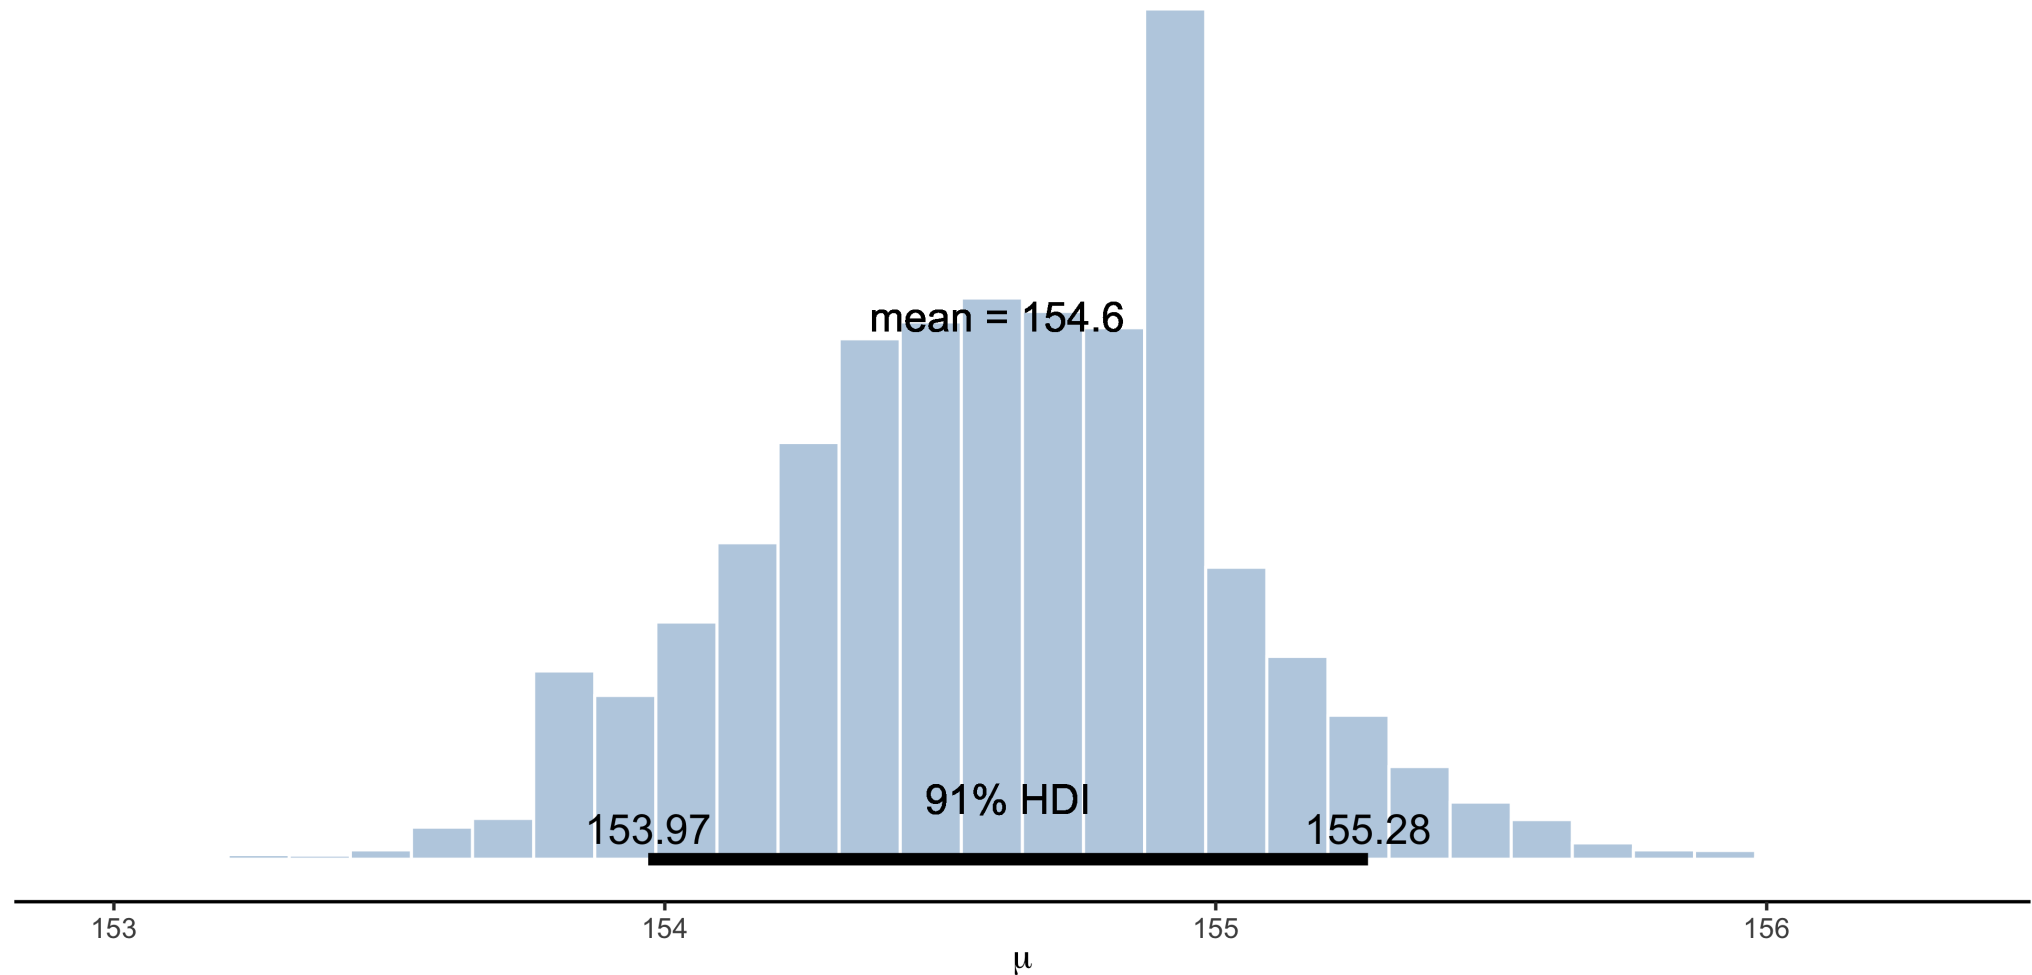
Posterior distribution - Grid approximation

```
1 sample.rows <- sample(x = 1:nrow(post), size = 1e4, replace = TRUE, prob = post$prob)
2 sample.mu <- post$mu[sample.rows]
3 sample.sigma <- post$sigma[sample.rows]
```




Posterior distribution - Marginal distributions

```
1 posterior_plot(samples = sample.mu, nbins = 30) + labs(x = expression(mu))
```



Posterior distribution - Marginal distributions

```
1 posterior_plot(samples = sample.sigma, nbins = 30) + labs(x = expression(sigma))
```



Introduction to brms



Introduction to brms

Under the hood: **Stan** is a probabilistic programming language written in **C++**, which implements several MCMC algorithms: HMC, NUTS, L-BFGS...

```
1 data {
2   int<lower=0> J; // number of schools
3   real y[J]; // estimated treatment effects
4   real<lower=0> sigma[J]; // s.e. of effect estimates
5 }
6
7 parameters {
8   real mu;
9   real<lower=0> tau;
10  real eta[J];
11 }
12
13 transformed parameters {
14   real theta[J];
15   for (j in 1:J)
16     theta[j] = mu + tau * eta[j];
17 }
18
19 model {
20   target += normal_lpdf(eta | 0, 1);
21   target += normal_lpdf(y | theta, sigma);
22 }
```

Bayesian regression models using Stan

The `brms` package ([Bürkner, 2017](#)) can be used to fit multilevel (or single-level) linear (or not) Bayesian regression models in `Stan` but using the intuitive syntax of `lme4` (cf. our tutorial paper, [Nalborczyk et al., 2019](#)).

For instance, the following model:

$$y_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \alpha_{\text{subject}[i]} + \alpha_{\text{item}[i]} + \beta x_i$$

is specified with `brms` (as with `lme4`) as follows:

```
1 model <- brm(y ~ x + (1 | subject) + (1 | item), data = d, family = gaussian() )
```



Syntax reminders

The **brms** package uses the same syntax as the base **R** functions (such as **lm**) or functions from the **lme4** package.

```
1 Reaction ~ Days + (1 + Days | Subject)
```

The left-hand side represents our dependent variable (or outcome, that is, what we are trying to predict). The **brms** package can also be used to fit multivariate models (several outcomes) by combining them with **mvbind()**.

```
1 mvbind(Reaction, Memory) ~ Days + (1 + Days | Subject)
```

The right-hand side is used to define the predictors. The intercept is generally implicit, so that the two formulations below are equivalent.

```
1 mvbind(Reaction, Memory) ~ Days + (1 + Days | Subject)
2 mvbind(Reaction, Memory) ~ 1 + Days + (1 + Days | Subject)
```



Syntax reminders

If you want to fit a model without an intercept (because why not), you must specify it explicitly, as shown below.

```
1 mvbind(Reaction, Memory) ~ 0 + Days + (1 + Days | Subject)
```

By default `brms` assumes a Gaussian likelihood function. This assumption can be modified easily by specifying the desired likelihood via the `family` argument.

```
1 brm(Reaction ~ 1 + Days + (1 + Days | Subject), family = lognormal() )
```

In case of doubt, read the documentation (it's very exciting to read) available at [?brm](#).



Some useful functions

```
1 # retrieving the Stan code generated by brms
2 make_stancode(formula, ...)
3 stancode(fit)
4
5 # checking and/or defining priors
6 get_prior(formula, ...)
7 set_prior(prior, ...)
8
9 # retrieving predictions from a brms model
10 fitted(fit, ...)
11 predict(fit, ...)
12 conditional_effects(fit, ...)
13
14 # posterior predictive checking
15 pp_check(fit, ...)
16
17 # model comparison and hypothesis testing
18 loo(fit1, fit2, ...)
19 bayes_factor(fit1, fit2, ...)
20 model_weights(fit1, fit2, ...)
21 hypothesis(fit, hypothesis, ...)
```



A first example

```
1 library(brms)
2 mod1 <- brm(height ~ 1, data = d2)
```

```
1 posterior_summary(mod1, pars = c("^b_", "sigma"), probs = c(0.025, 0.975) )
```

	Estimate	Est.Error	Q2.5	Q97.5
b_Intercept	154.614878	0.4224360	153.787631	155.452195
sigma	7.763965	0.2936431	7.217446	8.354354

These data represent the marginal (posterior) distributions of each parameter. In other words, the posterior probability μ , averaged over all possible values of σ , is best described by a Gaussian distribution with a mean of 154.6 and a standard deviation of 0.42. The credibility interval (\neq confidence interval) indicates the 95% most plausible values of μ or σ (given the data and priors).

Using our prior

By default `brms` uses a very uninformative prior centred on the mean value of the measured variable. We can therefore refine the estimate using our knowledge of the usual distribution of heights in humans.



The `get_prior()` function is used to display a list of default priors as well as all the priors we can specify, given a certain model formula (i.e., a way of writing our model) and a set of data.

```
1 get_prior(height ~ 1, data = d2)
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	source
student_t(3, 154.3, 8.5)	Intercept									default
student_t(3, 0, 8.5)	sigma							0		default

Using our prior

```
1 priors <- c(  
2   prior(normal(174, 20), class = Intercept),  
3   prior(exponential(0.01), class = sigma)  
4 )  
5  
6 mod2 <- brm(  
7   height ~ 1,  
8   prior = priors,  
9   family = gaussian(),  
10  data = d2  
11 )
```



Using our prior

```
1 summary(mod2)
```

```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: height ~ 1
Data: d2 (Number of observations: 352)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000

Population-Level Effects:
      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
Intercept  154.60      0.41  153.80  155.42 1.00    3433    2656



Family Specific Parameters:
      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
sigma      7.76      0.30   7.23   8.39 1.00    3032    2499

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
and Tail_ESS are effective sample size measures, and Rhat is the potential
scale reduction factor on split chains (at convergence, Rhat = 1).
```



Using a more informative prior

```
1 priors <- c(  
2   prior(normal(174, 0.1), class = Intercept),  
3   prior(exponential(0.01), class = sigma)  
4 )  
5  
6 mod3 <- brm(  
7   height ~ 1,  
8   prior = priors,  
9   family = gaussian(),  
10  data = d2  
11 )
```



Using a more informative prior

```
1 summary(mod3)
```

```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: height ~ 1
Data: d2 (Number of observations: 352)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000

Population-Level Effects:
      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
Intercept    173.84      0.10   173.65   174.03 1.00     3108     2389

Family Specific Parameters:
      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
sigma     20.81      0.79    19.33    22.48 1.00     2908     2312

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
and Tail_ESS are effective sample size measures, and Rhat is the potential
scale reduction factor on split chains (at convergence, Rhat = 1).
```

We note that the estimated value for μ has hardly “moved” from the prior...but we can also see that the estimated value for σ has greatly increased. What happened? We told the model that we were fairly certain of our μ value, the model then “adapted”, which explains the large value of σ ...



Prior precision (heuristic)

Prior distributions can generally be considered as posterior distributions obtained from previous data.

We know that the σ of a Gaussian posterior is given by:

$$\sigma_{\text{post}} = 1 / \sqrt{n}$$


Which implies a **quantity of data** $n = 1 / \sigma_{\text{post}}^2$. Our prior had a $\sigma = 0.1$, which implies $n = 1 / 0.1^2 = 100$.

We can therefore consider that the prior $\text{Normal}(174, 0.1)$ is equivalent to the case in which we would have observed 100 heights with mean 174.



Visualising samples from the posterior distribution

```
1 post <- as_draws_df(x = mod2) %>%  
2   mutate(density = get_density(b_Intercept, sigma, n = 1e2) )  
3  
4 ggplot(post, aes(x = b_Intercept, y = sigma, color = density) ) +  
5   geom_point(size = 2, alpha = 0.5, show.legend = FALSE) +  
6   labs(x = expression(mu), y = expression(sigma) ) +  
7   viridis::scale_color_viridis()
```



Retrieving samples from the posterior distribution

```
1 # gets the first 6 samples
2 head(post)
```


```
# A draws_df: 6 iterations, 1 chains, and 5 variables
  b_Intercept sigma lprior  lp__ density
1         154    8.2   -9.1 -1229  0.082
2         155    7.4   -9.0 -1229  0.150
3         155    7.2   -9.0 -1229  0.139
4         154    8.0   -9.1 -1227  0.582
5         155    8.3   -9.1 -1228  0.263
6         155    7.6   -9.0 -1228  0.215
# ... hidden reserved variables {'.chain', '.iteration', '.draw'}
```

```
1 # gets the median and the 95% credible interval
2 t(sapply(post[, 1:2], quantile, probs = c(0.025, 0.5, 0.975) ) )
```


	2.5%	50%	97.5%
b_Intercept	153.795472	154.59952	155.418792
sigma	7.227358	7.74985	8.386699

Visualising the posterior distribution

```
1 H.scv <- Hscv(post[, 1:2])
2 fhat_post <- kde(x = post[, 1:2], H = H.scv, compute.cont = TRUE)
3
4 plot(fhat_post, display = "persp", col = "purple", border = NA,
5      xlab = "\nmu", ylab = "\nsigma", zlab = "\np(mu, sigma)",
6      shade = 0.8, phi = 30, ticktype = "detailed",
7      cex.lab = 1.2, family = "Helvetica")
```




Visualising the posterior distribution



Including predictors

How does height change with weight?

```
1 d2 %>%  
2   ggplot(aes(x = weight, y = height)) +  
3   geom_point(colour = "white", fill = "black", pch = 21, size = 3, alpha = 0.8)
```




Linear regression with a single continuous predictor

$$h_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta x_i$$

```
1 linear_model <- lm(height ~ weight, data = d2)
2 rethinking::precis(linear_model, prob = 0.95)
```

	mean	sd	2.5%	97.5%
(Intercept)	113.8793936	1.91106523	110.1337746	117.6250126
weight	0.9050291	0.04204752	0.8226175	0.9874407



Notations

Let's consider a linear regression model with a single predictor, a slope, an intercept, and residuals distributed according to a normal distribution. The notation:

$$h_i = \alpha + \beta x_i + \epsilon_i \quad \text{avec} \quad \epsilon_i \sim \text{Normal}(0, \sigma)$$

is equivalent to:

$$h_i - (\alpha + \beta x_i) \sim \text{Normal}(0, \sigma)$$

rearranging the expression:

$$h_i \sim \text{Normal}(\alpha + \beta x_i, \sigma).$$

The above notations are equivalent, but the last one is more flexible, and will allow us to extend it more intuitively to multilevel models.



Linear regression with a single continuous predictor

$$h_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta x_i$$

$$\alpha \sim \text{Normal}(174, 20)$$

$$\beta \sim \text{Normal}(0, 10)$$

$$\sigma \sim \text{Exponential}(0.01)$$

In this model, μ is no longer a parameter to be estimated (because μ is **determined** by α and β). Instead, we will estimate α and β .

Reminders: α is the intercept, that is, the expected height, when the weight is equal to 0. β is the slope, that is, the expected change in height when weight increases by one unit.



Linear regression with a single continuous predictor

```
1 priors <- c(  
2   prior(normal(174, 20), class = Intercept),  
3   prior(normal(0, 10), class = b),  
4   prior(exponential(0.01), class = sigma)  
5 )  
6  
7 mod4 <- brm(  
8   height ~ 1 + weight,  
9   prior = priors,  
10  family = gaussian(),  
11  data = d2  
12 )
```



Linear regression with a single continuous predictor

```
1 posterior_summary(mod4)
```

	Estimate	Est.Error	Q2.5	Q97.5
b_Intercept	113.9022663	1.89453254	110.2751443	117.5863999
b_weight	0.9045325	0.04172372	0.8226575	0.9834808
sigma	5.1023966	0.19607421	4.7309481	5.4967173
lprior	-12.2671372	0.01272490	-12.2925817	-12.2432602
lp__	-1083.1151428	1.14546650	-1086.1172584	-1081.7573733

- $\alpha = 113.90$, 95% CrI [110.28, 117.59] represents the expected height when weight equals 0kg...
- $\beta = 0.90$, 95% CrI [0.82, 0.98] indicates that for an increase of 1kg, we can expect an increase of 0.90cm.

Linear regression with a single continuous predictor

```
1 d2$weight.c <- d2$weight - mean(d2$weight)
2
3 mod5 <- brm(
4   height ~ 1 + weight.c,
5   prior = priors,
6   family = gaussian(),
7   data = d2
8 )
```


```
1 fixef(mod5) # retrieving the fixed effects estimates
```

	Estimate	Est.Error	Q2.5	Q97.5
Intercept	154.5892334	0.26511032	154.0580097	155.1021441
weight.c	0.9048262	0.04247021	0.8194467	0.9874537

If the predictor (weight) is centred, the intercept represents the expected value of height (in cm) when weight is at its mean value.

Visualising predictions from the model

```
1 d2 %>%  
2   ggplot(aes(x = weight, y = height) ) +  
3   geom_point(colour = "white", fill = "black", pch = 21, size = 3, alpha = 0.8) +  
4   geom_abline(intercept = fixef(mod4)[1], slope = fixef(mod4)[2], lwd = 1)
```



Visualising the uncertainty on μ using fitted()

```


1 # defining a grid of possible values for "weight"
2 weight.seq <- data.frame(weight = seq(from = 25, to = 70, by = 1) )
3
4 # retrieving the model's predicted mus (average height) for these values of "weight"
5 mu <- data.frame(fitted(mod4, newdata = weight.seq) ) %>% bind_cols(weight.seq)
6
7 # displaying the first 10 rows of mu
8 head(mu, 10)

```

	Estimate	Est.Error	Q2.5	Q97.5	weight
1	136.5156	0.8729446	134.8388	138.2282	25
2	137.4201	0.8331902	135.8118	139.0699	26
3	138.3246	0.7936381	136.7802	139.8924	27
4	139.2292	0.7543200	137.7739	140.7197	28
5	140.1337	0.7152745	138.7559	141.5451	29
6	141.0382	0.6765488	139.7387	142.3767	30
7	141.9428	0.6382011	140.7096	143.1942	31
8	142.8473	0.6003039	141.6932	144.0216	32
9	143.7518	0.5629482	142.6682	144.8488	33
10	144.6564	0.5262493	143.6445	145.6845	34

Visualising the uncertainty on μ using fitted()

```
1 d2 %>%  
2   ggplot(aes(x = weight, y = height) ) +  
3   geom_point(colour = "white", fill = "black", pch = 21, size = 3, alpha = 0.8) +  
4   geom_smooth(  
5     data = mu, aes(y = Estimate, ymin = Q2.5, ymax = Q97.5),  
6     stat = "identity",  
7     color = "black", alpha = 0.8, size = 1  
8   )
```



Prediction intervals (incorporating σ)


As a reminder, our model is defined as: $h_i \sim \text{Normal}(\alpha + \beta x_i, \sigma)$. For the moment, we have only represented the predictions for μ . How can we incorporate σ into our predictions?

```
1 # defining a grid of possible values for "weight"
2 weight.seq <- data.frame(weight = seq(from = 25, to = 70, by = 1) )
3
4 # retrieving the model's predicted heights for these values of "weight"
5 pred_height <- data.frame(predict(mod4, newdata = weight.seq) ) %>% bind_cols(weight.seq)
6
7 # displaying the first 10 rows of pred_height
8 head(pred_height, 10)
```

	Estimate	Est.Error	Q2.5	Q97.5	weight
1	136.6003	5.241255	126.6961	146.9625	25
2	137.5269	5.123556	127.5401	147.6144	26
3	138.2405	5.176174	128.1914	148.3255	27
4	139.3099	5.243116	129.3009	149.4473	28
5	140.0189	5.168394	129.7397	150.1358	29
6	140.8688	5.218597	130.8407	151.0752	30
7	142.0739	5.152610	132.1567	152.0562	31
8	142.7756	5.218983	132.5639	153.1091	32
9	143.7965	5.244761	133.4364	154.0000	33
10	144.6657	5.198395	134.3009	154.5210	34

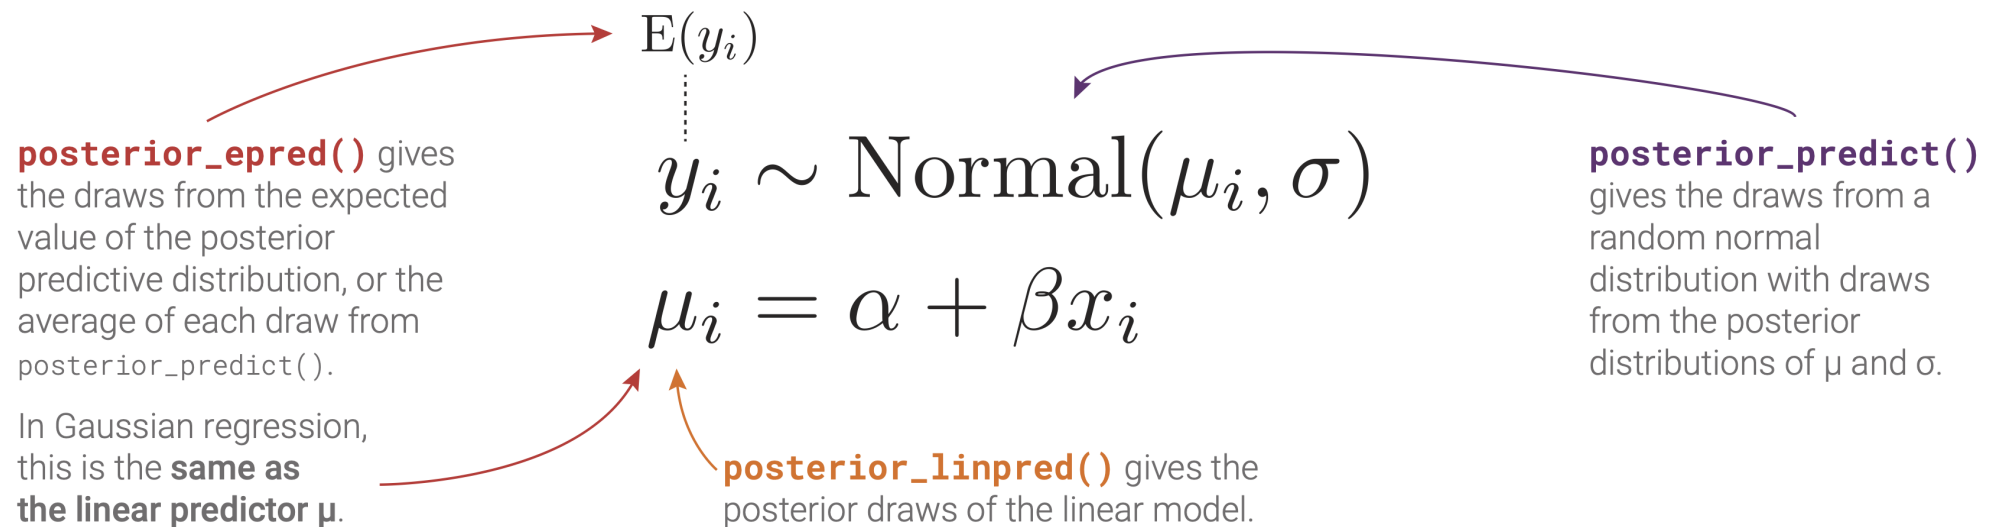
Prediction intervals (incorporating σ)

```
1 d2 %>%  
2   ggplot(aes(x = weight, y = height) ) +  
3   geom_point(colour = "white", fill = "black", pch = 21, size = 3, alpha = 0.8) +  
4   geom_ribbon(  
5     data = pred_height, aes(x = weight, ymin = Q2.5, ymax = Q97.5),  
6     alpha = 0.2, inherit.aes = FALSE  
7   ) +  
8   geom_smooth(  
9     data = mu, aes(y = Estimate, ymin = Q2.5, ymax = Q97.5),  
10    stat = "identity", color = "black", alpha = 0.8, size = 1  
11  )
```



Built-in brms functions

The `brms` package also includes the `posterior_epred()`, `posterior_linpred()`, and `posterior_predict()` functions, which can be used to generate predictions from models fitted with `brms`. Andrew Heiss gives a detailed description of how these functions work in [this blog post](#).



Reminder: Two types of uncertainty

Two sources of uncertainty in the model: uncertainty concerning the estimation of the value of the parameters but also uncertainty concerning the sampling process.

Epistemic uncertainty: The posterior distribution orders all possible combinations of parameter values according to their relative plausibility.


Aleatory uncertainty : The distribution of simulated data is a distribution that contains uncertainty related to a sampling process (i.e., generating data from a Gaussian distribution).

See also this short article by O'Hagan ([2004](#)).



Polynomial regression

```
1 d %>% # note that we use d instead of d2
2   ggplot(aes(x = weight, y = height)) +
3   geom_point(colour = "white", fill = "black", pch = 21, size = 3, alpha = 0.8)
```




Standardised (scaled) scores

```
1 d <- d %>% mutate(weight.s = (weight - mean(weight) ) / sd(weight) )  
2  
3 d %>%  
4   ggplot(aes(x = weight.s, y = height) ) +  
5   geom_point(colour = "white", fill = "black", pch = 21, size = 3, alpha = 0.8)
```

```
1 c(mean(d$weight.s), sd(d$weight.s) )
```

```
[1] -2.712698e-18  1.000000e+00
```



Standardised (scaled) scores

Why standardising predictors?

- **Interpretation.** Easier to compare the coefficients of several predictors. A change of one standard deviation in the predictor corresponds to a change of one standard deviation in the response (if the response is also standardised).
- **Fitting.** When the predictors contain large values (or values that are too different from one another), this can lead to convergence problems (cf. Course n°03).



Polynomial regression - Exercise

$$h_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_1 x_i + \beta_2 x_i^2$$

$$\alpha \sim \text{Normal}(156, 100)$$

$$\beta_1, \beta_2 \sim \text{Normal}(0, 10)$$

$$\sigma \sim \text{Exponential}(0.01)$$

It's up to you to build and fit this model using `brms::brm()`.



Polynomial regression

```
1 priors <- c(  
2   prior(normal(156, 100), class = Intercept),  
3   prior(normal(0, 10), class = b),  
4   prior(exponential(0.01), class = sigma)  
5 )  
6  
7 mod6 <- brm(  
8   # NB: polynomials should be written with the I() function...  
9   height ~ 1 + weight.s + I(weight.s^2),  
10  prior = priors,  
11  family = gaussian(),  
12  data = d  
13 )
```



Polynomial regression

```
1 summary(mod6)
```

```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: height ~ 1 + weight.s + I(weight.s^2)
Data: d (Number of observations: 544)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

Population-Level Effects:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	146.65	0.37	145.95	147.39	1.00	3438	2694
weight.s	21.40	0.29	20.84	21.99	1.00	3198	2387
Iweight.sE2	-8.41	0.28	-8.96	-7.85	1.00	3117	2845

Family Specific Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	5.78	0.17	5.46	6.15	1.00	3352	3090

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).



Visualising the model's predictions

```


1 # defining a grid of possible (standardised) values for "weight"
2 weight.seq <- data.frame(weight.s = seq(from = -2.5, to = 2.5, length.out = 50) )
3
4 # retrieving the model's predictions for these values of weight
5 mu <- data.frame(fitted(mod6, newdata = weight.seq) ) %>% bind_cols(weight.seq)
6 pred_height <- data.frame(predict(mod6, newdata = weight.seq) ) %>% bind_cols(weight.seq)
7
8 # displaying the first ten rows of pred_height
9 head(pred_height, 10)

```

	Estimate	Est.Error	Q2.5	Q97.5	weight.s
1	40.60063	6.010429	28.97013	52.51992	-2.500000
2	47.10254	5.955673	35.32591	58.77250	-2.397959
3	53.19186	6.006070	41.31088	65.14786	-2.295918
4	59.29052	5.760499	47.61468	70.36520	-2.193878
5	65.32630	5.963502	53.34686	76.94875	-2.091837
6	70.76771	5.769429	59.88322	82.18180	-1.989796
7	76.26657	5.815825	64.92816	87.47296	-1.887755
8	81.64095	5.787103	70.60726	92.95473	-1.785714
9	86.67635	5.665041	75.79001	97.45832	-1.683673
10	91.71143	5.880964	80.16754	103.33322	-1.581633

Visualising the model's predictions

```
1 d %>%  
2   ggplot(aes(x = weight.s, y = height) ) +  
3   geom_point(colour = "white", fill = "black", pch = 21, size = 3, alpha = 0.8) +  
4   geom_ribbon(  
5     data = pred_height, aes(x = weight.s, ymin = Q2.5, ymax = Q97.5),  
6     alpha = 0.2, inherit.aes = FALSE  
7   ) +  
8   geom_smooth(  
9     data = mu, aes(y = Estimate, ymin = Q2.5, ymax = Q97.5),  
10    stat = "identity", color = "black", alpha = 0.8, size = 1  
11  )
```



Effect sizes in regression models

There are several methods for computing effect sizes in Bayesian models. Gelman & Pardoe ([2006](#)) propose a method for computing a sample-based R^2 .

Marsman & Wagenmakers ([2017](#)) and Marsman et al. ([2019](#)) generalise existing methods to compute a ρ^2 for ANOVA designs (i.e., with categorical predictors), which represents an estimate of the effect size in the population.

“

Similar to most of the ES measures that have been proposed for the ANOVA model, the squared multiple correlation coefficient ρ^2 [...] is a so-called proportional reduction in error measure (PRE). In general, a PRE measure expresses the proportion of the variance in an outcome y that is attributed to the independent variables x ([Marsman et al., 2019](#)).



Effect sizes in regression models

$$\rho^2 = \frac{\sum_{i=1}^n \pi_i (\beta_i - \beta)^2}{\sigma^2 + \sum_{i=1}^n \pi_i (\beta_i - \beta)^2}$$

$$\rho^2 = \frac{\frac{1}{n} \sum_{i=1}^n \beta_i^2}{\sigma^2 + \frac{1}{n} \sum_{i=1}^n \beta_i^2}$$


$$\rho^2 = \frac{\beta^2 \tau^2}{\sigma^2 + \beta^2 \tau^2}$$

```
1 post <- as_draws_df(x = mod4)
2 beta <- post$b_weight
3 sigma <- post$sigma
4 rho <- beta^2 * var(d2$weight) / (sigma^2 + beta^2 * var(d2$weight) )
```

Caution: if there are several predictors, it depends on the covariance structure...

Effect sizes in regression models

```
1 posterior_plot(samples = rho, usemode = TRUE) + labs(x = expression(rho) )
```



```
1 summary(lm(height ~ weight, data = d2) )$r.squared
```

```
[1] 0.5696444
```




Effect sizes in regression models

```
1 bayes_R2(mod4)
```

	Estimate	Est.Error	Q2.5	Q97.5
R2	0.568025	0.02283433	0.51965	0.607488

```
1 bayes_R2(mod4, summary = FALSE)[, 1] %>%  
2   posterior_plot(usemode = TRUE) +  
3   labs(x = expression(rho) )
```



Summary

A new model with two and then three parameters was presented: the Gaussian model, then Gaussian linear regression model, making it possible to relate two continuous variables.

As previously, Bayes' theorem is used to update our prior knowledge of parameter values into posterior knowledge, a synthesis between our priors and the information contained in the data.

The `brms` package can be used to fit all sorts of models with a syntax similar to that used by `lm()`.

The `fitted()` function is used to retrieve the predictions of a model.

The `predict()` function is used to simulate data from a model fitted with `brms`.



Practical work - 1/2

Select all rows in the `howell1` dataset corresponding to minors individuals (age < 18). This should result in a dataframe of 192 rows.

Fit a linear regression model using the `brms::brm()` function. Report and interpret the estimates from this model. For an increase of 10 units in `weight`, what increase in height (`height`) does the model predict?

Plot the raw data with `weight` on the x-axis and `height` on the y-axis. Overlay the model's regression line of the model and an 89% credibility interval for the mean. Add an 89% credibility interval for the predicted sizes.

What do you think of the 'fit' of the model? What assumptions of the model would you be willing to modify in order to improve the model's fit?



Practical work - 2/2

Let's say you've consulted a colleague who's an expert in [allometry](#) (i.e., the phenomena of differential organ growth) and she explains to you that it doesn't make sense to model the relationship between weight and height... because we know that it's the logarithm of weight which is (linearly) related to height!

Model the relationship between height (cm) and log weight (log-kg). Use the entire `howe11` dataframe (all 544 rows). Fit the following model using `brms::brm()`.

$$h_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta \cdot \log(w_i)$$

$$\alpha \sim \text{Normal}(174, 100)$$

$$\beta \sim \text{Normal}(0, 100)$$

$$\sigma \sim \text{Exponential}(0.01)$$

Where h_i is the height of individual i and w_i the weight of individual i . The function for calculating the log in **R** is simply `log()`. Can you interpret the results? Hint: plot the raw data and superimpose the model's predictions...



Proposed solution

```
1 # we keep only the individuals younger than 18
2 d <- open_data(howell) %>% filter(age < 18)
3
4 priors <- c(
5   prior(normal(150, 100), class = Intercept),
6   prior(normal(0, 10), class = b),
7   prior(exponential(0.01), class = sigma)
8 )
9
10 mod7 <- brm(
11   height ~ 1 + weight,
12   prior = priors,
13   family = gaussian(),
14   data = d
15 )
```



Proposed solution

```
1 summary(mod7, prob = 0.89)
```



```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: height ~ 1 + weight
Data: d (Number of observations: 192)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000

Population-Level Effects:
      Estimate Est.Error l-89% CI u-89% CI Rhat Bulk_ESS Tail_ESS
Intercept    58.22     1.40   56.00   60.47 1.00    3853    3011
weight        2.72     0.07    2.61    2.83 1.00    3445    2962

Family Specific Parameters:
      Estimate Est.Error l-89% CI u-89% CI Rhat Bulk_ESS Tail_ESS
sigma        8.53     0.45    7.84    9.27 1.00    4283    2872

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
and Tail_ESS are effective sample size measures, and Rhat is the potential
scale reduction factor on split chains (at convergence, Rhat = 1).
```



Proposed solution

```

1 # defining a grid of possible values for "weight"
2 weight.seq <- data.frame(weight = seq(from = 5, to = 45, length.out = 1e2) )
3
4 # retrieving the model's predictions for these values fo "weight"
5 mu <- data.frame(
6   fitted(mod7, newdata = weight.seq, probs = c(0.055, 0.945) )
7 ) %>%
8   bind_cols(weight.seq)
9
10 pred_height <- data.frame(
11   predict(mod7, newdata = weight.seq, probs = c(0.055, 0.945) )
12 ) %>%
13   bind_cols(weight.seq)
14
15 # displaying the first 6 rows of pred_height
16 head(pred_height)

```


	Estimate	Est.Error	Q5.5	Q94.5	weight
1	71.90158	8.558478	58.38327	85.76683	5.000000
2	72.80918	8.809837	58.64175	86.74542	5.404040
3	74.07137	8.558394	60.51732	87.58434	5.808081
4	75.08137	8.717473	61.08438	89.05932	6.212121
5	75.95542	8.633522	62.39603	89.61429	6.616162
6	77.39934	8.498564	63.69033	91.05000	7.020202

Proposed solution

```

1 d %>%
2   ggplot(aes(x = weight, y = height) ) +
3   geom_point(colour = "white", fill = "black", pch = 21, size = 3, alpha = 0.8) +
4   geom_ribbon(
5     data = pred_height, aes(x = weight, ymin = Q5.5, ymax = Q94.5),
6     alpha = 0.2, inherit.aes = FALSE
7   ) +
8   geom_smooth(
9     data = mu, aes(y = Estimate, ymin = Q5.5, ymax = Q94.5),
10    stat = "identity", color = "black", alpha = 0.8, size = 1
11  )

```



Proposed solution

```
1 # we now consider all individuals
2 d <- open_data(howell)
3
4 mod8 <- brm(
5   # we now predict height using the logarithm of weight
6   height ~ 1 + log(weight),
7   prior = priors,
8   family = gaussian(),
9   data = d
10  )
```



Proposed solution

```
1 summary(mod8, prob = 0.89)
```



```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: height ~ 1 + log(weight)
Data: d (Number of observations: 544)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000

Population-Level Effects:
      Estimate Est.Error l-89% CI u-89% CI Rhat Bulk_ESS Tail_ESS
Intercept   -23.56      1.34  -25.67  -21.42 1.00    4356    3128
logweight    47.01      0.38   46.39   47.63 1.00    4396    2952

Family Specific Parameters:
      Estimate Est.Error l-89% CI u-89% CI Rhat Bulk_ESS Tail_ESS
sigma      5.15      0.15    4.91    5.41 1.00    4641    2770

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
and Tail_ESS are effective sample size measures, and Rhat is the potential
scale reduction factor on split chains (at convergence, Rhat = 1).
```



Proposed solution

```

1 # defining a grid of possible values for "weight"
2 weight.seq <- data.frame(weight = seq(from = 5, to = 65, length.out = 1e2) )
3
4 # retrieving the model's predictions for these values of "weight"
5 mu <- data.frame(
6   fitted(mod8, newdata = weight.seq, probs = c(0.055, 0.945) )
7 ) %>%
8   bind_cols(weight.seq)
9
10 pred_height <- data.frame(
11   predict(mod8, newdata = weight.seq, probs = c(0.055, 0.945) )
12 ) %>%
13   bind_cols(weight.seq)
14
15 # displaying the first 6 rows of "pred_height"
16 head(pred_height)


```

	Estimate	Est.Error	Q5.5	Q94.5	weight
1	52.08114	5.364763	43.64659	60.73506	5.000000
2	57.47366	5.180466	49.20607	65.72344	5.606061
3	62.29115	5.289523	53.91334	70.72491	6.212121
4	66.70296	5.104535	58.45216	74.90719	6.818182
5	70.49918	5.249748	62.17113	78.91841	7.424242
6	74.31789	5.177680	66.23217	82.75413	8.030303



Proposed solution

```
1 d %>%  
2   ggplot(aes(x = weight, y = height) ) +  
3   geom_point(colour = "white", fill = "black", pch = 21, size = 3, alpha = 0.8) +  
4   geom_ribbon(  
5     data = pred_height, aes(x = weight, ymin = Q5.5, ymax = Q94.5),  
6     alpha = 0.2, inherit.aes = FALSE  
7   ) +  
8   geom_smooth(  
9     data = mu, aes(y = Estimate, ymin = Q5.5, ymax = Q94.5),  
10    stat = "identity", color = "black", alpha = 0.8, size = 1  
11  )
```



References

- Bürkner, P.-C. (2017). **brms** : An *R* Package for Bayesian Multilevel Models Using *Stan*. *Journal of Statistical Software*, 80(1). <https://doi.org/10.18637/jss.v080.i01>
- Gelman, A., & Pardoe, I. (2006). Bayesian measures of explained variance and pooling in multilevel (hierarchical) models. *Technometrics*, 48(2), 241–251. <https://doi.org/10.1198/004017005000000517>
- Marsman, M., & Wagenmakers, E.-J. (2017). Three Insights from a Bayesian Interpretation of the One-Sided *P* Value. *Educational and Psychological Measurement*, 77(3), 529–539. <https://doi.org/10.1177/0013164416669201>
- Marsman, M., Waldorp, L., Dablander, F., & Wagenmakers, E.-J. (2019). Bayesian estimation of explained variance in ANOVA designs. *Statistica Neerlandica*, 0(0), 1–22. <https://doi.org/10.1111/stan.12173>
- Nalborczyk, L., Batailler, C., Løevenbruck, H., Vilain, A., & Bürkner, P.-C. (2019). An Introduction to Bayesian Multilevel Models Using brms: A Case Study of Gender Effects on Vowel Variability in Standard Indonesian. *Journal of Speech, Language, and Hearing Research*, 62(5), 1225–1242. https://doi.org/10.1044/2018_jslhr-s-18-0006
- O'Hagan, T. (2004). Dicing with the unknown. *Significance*, 1(3), 132–133. <https://doi.org/10.1111/j.1740-9713.2004.00050.x>



